

Introduction to Data Analytics

Xin Gao


Xin.gao@kaust.edu.sa

July 29, 2022

SDU


Neurons

- A **neuron** is an electrically excitable cell (threshold switching unit) that processes and transmits information by electrical and chemical signaling
 - Dendrites, axon, synapses
- The cell body of a neuron frequently gives rise to multiple dendrites, but never to more than one axon, although the axon may branch hundreds of times before it terminates



Neurons



- **Dendrites:** filaments that arise from the cell body, often extending for hundreds of micrometers and branching multiple times, giving rise to a complex "dendritic tree"
- **Axon:** a special cellular filament that arises from the cell body at a site called the axon hillock and travels for a distance, as far as 1m in humans or even more in other species
- **Synapses:** send signals from the axon of one neuron to a dendrite of another



Neurons

- We are born with about 100 billion neurons
- Computers are at least 10^6 times faster in raw switching speed
- But the brain is faster and reliable at computationally intensive tasks, such as computer vision, speech recognition, etc
- The brain is also fault-tolerant, and exhibits graceful degradation with damage
 - A neuron may connect to as many as 100,000 other neurons
 - Even if you break 50% of the connections, the brain can still function properly
 - Very strong and robust connection construction

Artificial Neurons




Artificial Neural Networks

- f might be non-linear function
- X (vector of) continuous and/or discrete variables
- Y (vector of) continuous and/or discrete variables
- Represent f by **network of logistic units**
- Each unit is a logistic function

$$\text{unit output} = \frac{1}{1 + \exp(w_0 + \sum_i w_i x_i)}$$

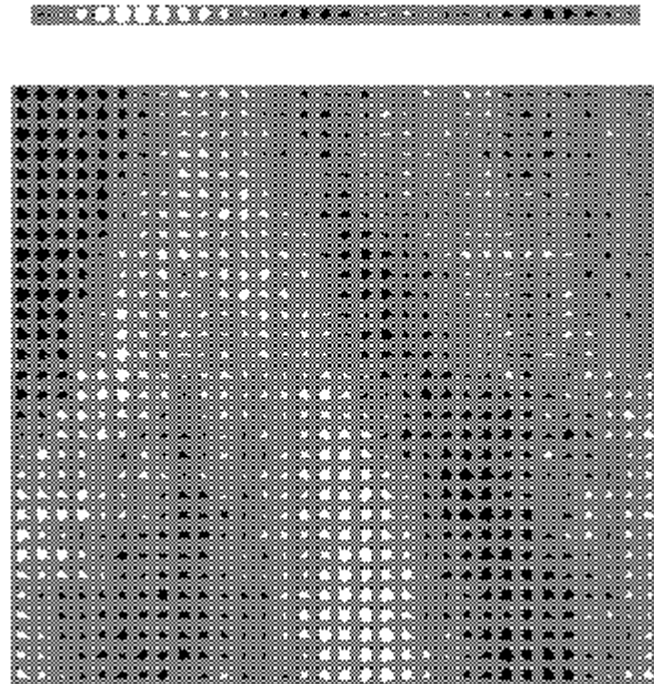

- Goal: train weights of all units to minimize the errors of predicted network outputs

Artificial Neural Networks



Example


- ALVINN: an autonomous land vehicle in a neural network – Pomerleau 1993



Connection Models

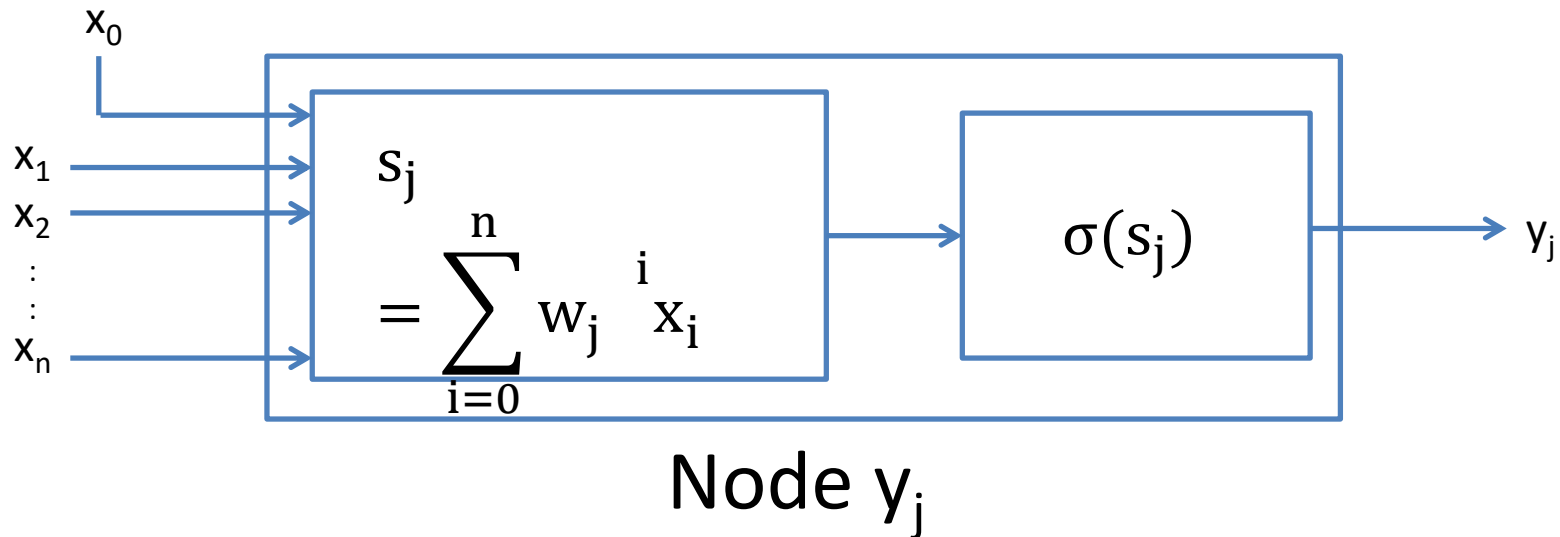
- Humans
 - Neuron switching time ~ 0.001 second
 - Number of neurons $\sim 10^{11}$
 - Connections per neuron $\sim 10^5$
 - Scene recognition time ~ 0.1 second
 - 100 inference steps doesn't seem enough
 - $>$ much parallel computation
- Properties of artificial neural networks (ANN's)
 - Many neuron-like threshold switching units
 - Many weighted interconnections among units
 - Highly parallel, distributed process

Artificial Neural Networks



- All nodes are involved in computation except for the input nodes, which simply send the input values to all nodes in the next layer


Artificial Neural Networks



- $s_j = \sum_{i=0}^n w_j^i x_i = w_j^0 x_0 + w_j^1 x_1 + \dots + w_j^n x_n$
- $x_0 = 1$, which is called "bias"
- $\sigma(s)$ is called the transfer function. There are various possibilities for $\sigma(s)$

Transfer Function


- Linear function
 - $\sigma(s) = ks$, where k is a real number
 - y_j is simply a linear function of the input x_i
 - A form of linear regression




Transfer Function

- Step function (threshold function)

$$-\sigma(s) = \begin{cases} 1, & s > 0 \\ -1, & s \leq 0 \end{cases}$$



$$-\sigma(s) = \begin{cases} 1, & s > 0 \\ 0, & s \leq 0 \end{cases}$$



Transfer Function


- Sigmoid function

- $\sigma(s) = \frac{1}{1+e^{-s}}$

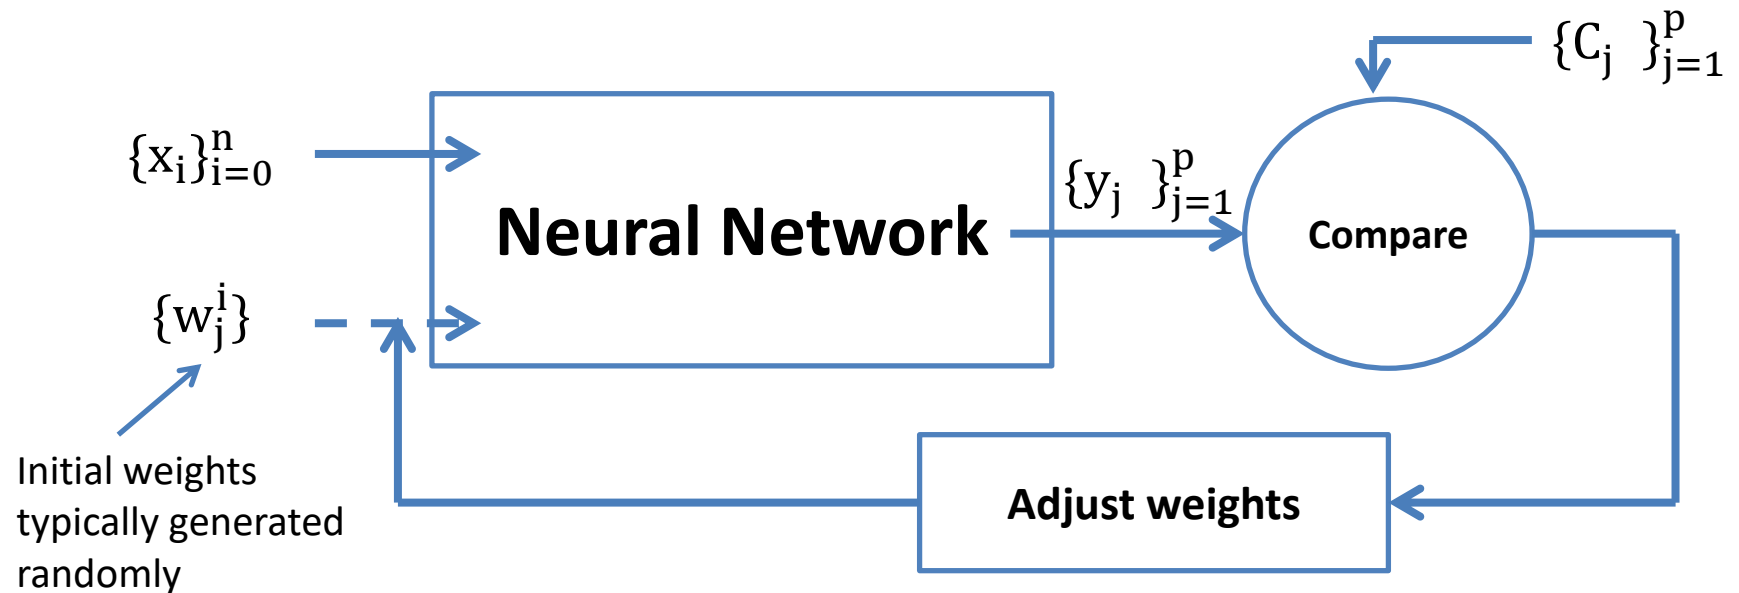
- Properties

- **Differentiable function:** a function whose derivative exists at each point in its domain

- $\sigma'(s) = \frac{d\sigma}{ds} = \frac{e^{-s}}{(1+e^{-s})^2} = \sigma(s)(1 - \sigma(s))$




Training ANN



- Training involves a “training set”, each member of the training set is a vector $\{x_i\}_{i=0}^n$ and an output $\{C_j\}_{j=1}^p$

Testing ANN



- Training involves a “test set”, each member of the test set is a vector $\{x_i\}_{i=0}^n$
- Training data and test data are separate data sets. However, they should be drawn from the same distribution

Training ANN

- We want to minimize the error
 - Least square error: $E = \frac{1}{2} \sum_{k=1}^p (y_k - C_k)^2$
- So E has to be minimized with respect to the weights $\{w_j^i\}$
- We need $\frac{\partial E}{\partial w_j^i}$ to discover how the error E depends on the $\{w_j^i\}$

Back-propagation Algorithm

- Gradient descent over entire network weight vector
- Will find a local, not necessarily global error minimum
 - In practice, often works well (can run multiple times)
- Minimizes error over training examples
 - Will it generalize well to subsequent examples?
- Training can take thousands of iterations. Slow!
- Using network after training is fast


Overfitting

- ANNs are supervised learning
 - Every supervised learning has risks of overfitting
- Training involves iterative weight updating. The number of iterations, n , is important
 - How do we choose n to minimize the error rate over future data?
 - We use cross validation

Expressive Capability of ANNs

- Boolean functions
 - Every boolean function can be represented by network with single hidden layer
 - But might require exponential (in number of inputs) hidden nodes
- Continuous functions
 - Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer [Cybenko 1989, Hornik et al. 1989]
 - Any function can be approximated to arbitrary accuracy by a network with two hidden layers [Cybenko 1988]

Example




Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001

Can this be learned?


Example

Learned hidden layer representation




Input		Hidden Values		Output
10000000	→	.89 .04 .08	→	10000000
01000000	→	.01 .11 .88	→	01000000
00100000	→	.01 .97 .27	→	00100000
00010000	→	.99 .97 .71	→	00010000
00001000	→	.03 .05 .02	→	00001000
00000100	→	.22 .99 .99	→	00000100
00000010	→	.80 .01 .98	→	00000010
00000001	→	.60 .94 .01	→	00000001


Example



Example




Example



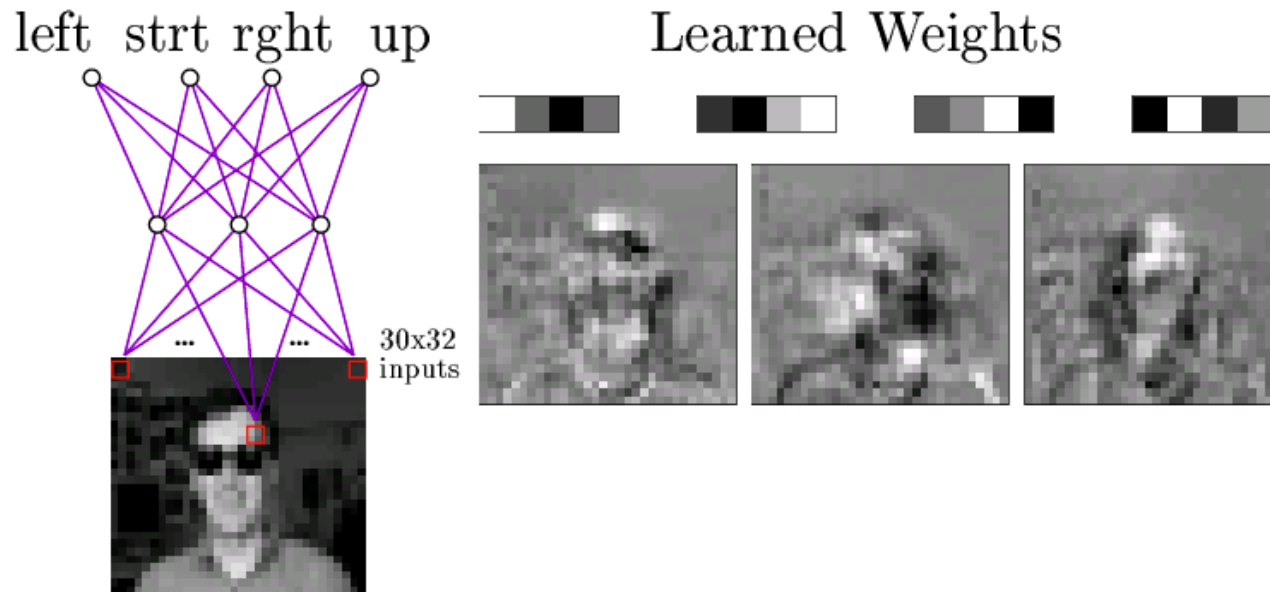
Another Example

- Neural network based face recognition



Typical input images


Another Example



Typical input images

Detour to Deep Neural Network

Deep neural network



What's the Problem?



618

¥509.00

德国马牌(Continental)轮胎/汽车
轮胎 285/35R16 91V CC6 本田
已卖500+条评价

德国马牌轮胎京东自营旗舰店
加入购物车



618

¥539.00

邓禄普(Dunlop)轮胎/汽车轮胎
225/65R17 102T ST30 CRV原配
已卖2200+条评价

京东自营专区
加入购物车



618

¥579.00

米其林(Michelin)轮胎/汽车轮胎
205/60R16 92V PRIMACY LC DT
已卖2500+条评价

米其林自营专区
加入购物车



618

¥529.00

普利司通(Bridgestone)轮胎/汽
车轮胎 205/55R16 91V ER300
已卖2000+条评价

京东自营专区
加入购物车



¥529.00

普利司通(Bridgestone)轮胎/汽
车轮胎 215/60R16 95V ER33UZ
已卖1200+条评价

京东自营专区
加入购物车



618

¥569.00

普利司通(Bridgestone)轮胎/汽
车轮胎 225/55R17 97W T001
已卖80+条评价

普利司通轮胎京东自营专区
加入购物车



618

¥618.00

米其林(Michelin)轮胎/汽车轮胎
215/60R16 99V PRIMACY3 ST
已卖1.6万+条评价

米其林自营专区
加入购物车



618

¥299.00

佳通(Giti)轮胎/汽车轮胎
205/55R16 91V 228 景城轮胎
已卖4000+条评价

京东自营专区
加入购物车



618

¥499.00

德国马牌(Continental)轮胎/汽
车轮胎 205/55R16 91V UC6 赛尔
已卖200+条评价

德国马牌轮胎京东自营旗舰店
加入购物车



618

¥439.00

米其林(Michelin)轮胎/汽车轮胎
195/65R15 91V PRIMACY3 ST
已卖1.6万+条评价

米其林自营专区
加入购物车

轮胎查查

What's the Problem?


 <p>¥509.00</p> <p>德国马牌 (Continental) 轮胎/汽车轮胎 205/55R16 91V CC6 本田 已有500+条评价</p> <p>德国马牌轮胎京东自营旗舰店</p> <p>加入购物车</p>	 <p>¥539.00</p> <p>邓禄普 (Dunlop) 轮胎/汽车轮胎 225/65R17 102T ST30 CRV/本田 已有2200+条评价</p> <p>京东自营专区</p> <p>加入购物车</p>	 <p>¥579.00</p> <p>米其林 (Michelin) 轮胎/汽车轮胎 205/60R16 92V PRIMACY LC DT 已有2500+条评价</p> <p>米其林自营专区</p> <p>加入购物车</p>	 <p>¥529.00</p> <p>普利司通 (Bridgestone) 轮胎/汽车轮胎 205/55R16 91V ER300 已有2000+条评价</p> <p>京东自营专区</p> <p>加入购物车</p>	 <p>¥529.00</p> <p>普利司通 (Bridgestone) 轮胎/汽车轮胎 215/60R16 95V ER33UZ 已有1200+条评价</p> <p>京东自营专区</p> <p>加入购物车</p>
 <p>¥569.00</p> <p>普利司通 (Bridgestone) 轮胎/汽车轮胎 225/55R17 97W T001 已有80+条评价</p> <p>普利司通轮胎京东自营专区</p> <p>加入购物车</p>	 <p>¥618.00</p> <p>米其林 (Michelin) 轮胎/汽车轮胎 215/60R16 99V PRIMACY3 ST 已有1.6万+条评价</p> <p>米其林自营专区</p> <p>加入购物车</p>	 <p>¥299.00</p> <p>佳通 (Giti) 轮胎/汽车轮胎 205/55R16 91V 228 奥城轮胎 已有4000+条评价</p> <p>京东自营专区</p> <p>加入购物车</p>	 <p>¥499.00</p> <p>德国马牌 (Continental) 轮胎/汽车轮胎 205/55R16 91V UC6 赛尔 已有200+条评价</p> <p>德国马牌轮胎京东自营旗舰店</p> <p>加入购物车</p>	 <p>¥439.00</p> <p>米其林 (Michelin) 轮胎/汽车轮胎 195/65R15 91V PRIMACY3 ST 已有1.6万+条评价</p> <p>米其林自营专区</p> <p>加入购物车</p>



Two Approaches in Supervised Learning

- Do we use the prediction performance to guide the search?
 - NO → **Filter**
 - Yes → **Wrapper**

Deep Learning – Convolutional Neural Network



LeNet-5


1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Convolution




Pooling

Convolutional Layer

Filters always extend the full depth of the input volume

32x32x3 image

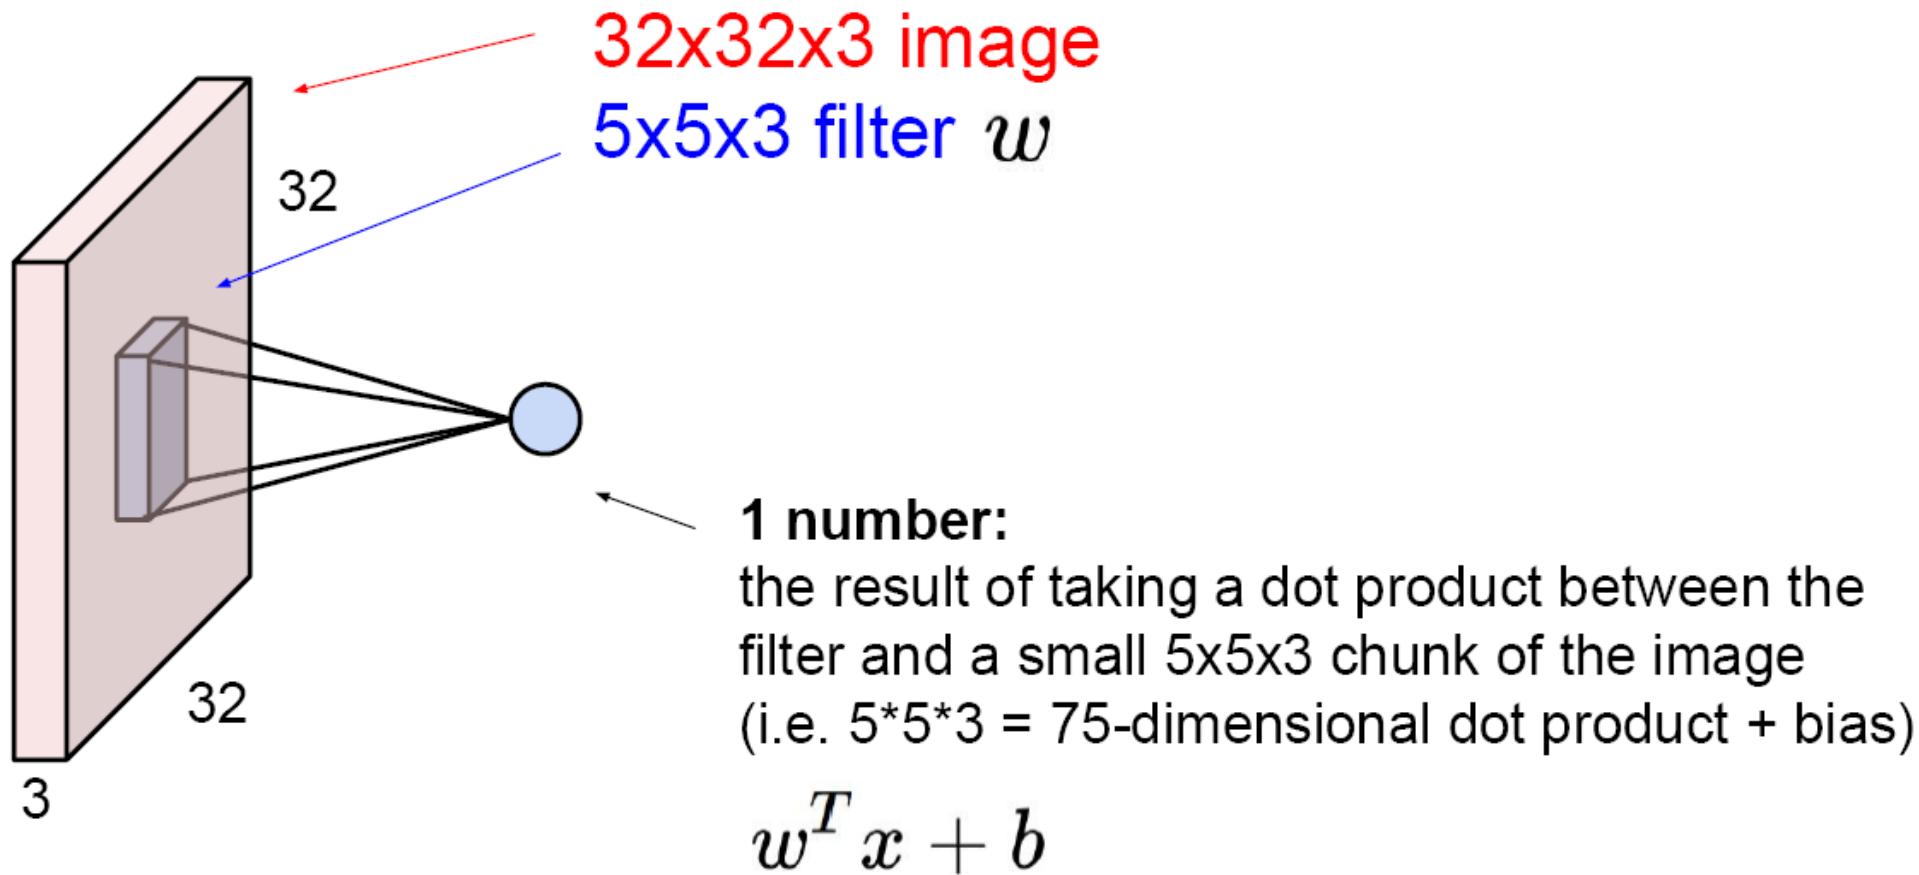


5x5x3 filter




Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”


Convolutional Layer



Convolutional Layer



Convolutional Layer



Convolutional Neural Network

Preview

[Zeiler and Fergus 2013]

Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].




Low-level features


Mid-level features

High-level features


Linearly separable classifier



VGG-16 Conv1_1




VGG-16 Conv3_2



VGG-16 Conv5_3


Convolutional Neural Network



Activation Function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Activation Function

$$\sigma(x) = 1/(1 + e^{-x})$$

- Squashes numbers to range [0,1]
- Historically popular since they have nice interpretation as a saturating “firing rate” of a neuron


3 problems:

1. Saturated neurons “kill” the gradients
2. Sigmoid outputs are not zero-centered
3. $\exp()$ is a bit compute expensive



Sigmoid

Activation Function




- Does not saturate (in +region)
- Very computationally efficient
- Converges much faster than sigmoid/tanh in practice (e.g. 6x)

ReLU
(Rectified Linear Unit)

- Not zero-centered output

Activation Function




- Does not saturate
- Computationally efficient
- Converges much faster than sigmoid/tanh in practice! (e.g. 6x)
- **will not “die”.**

Leaky ReLU

$$f(x) = \max(0.01x, x)$$


Pooling



Average Pooling

Max Pooling

Batch Normalization



1. compute the empirical mean and variance independently for each dimension.

2. Normalize

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

Softmax

Want to interpret raw classifier scores as **probabilities**

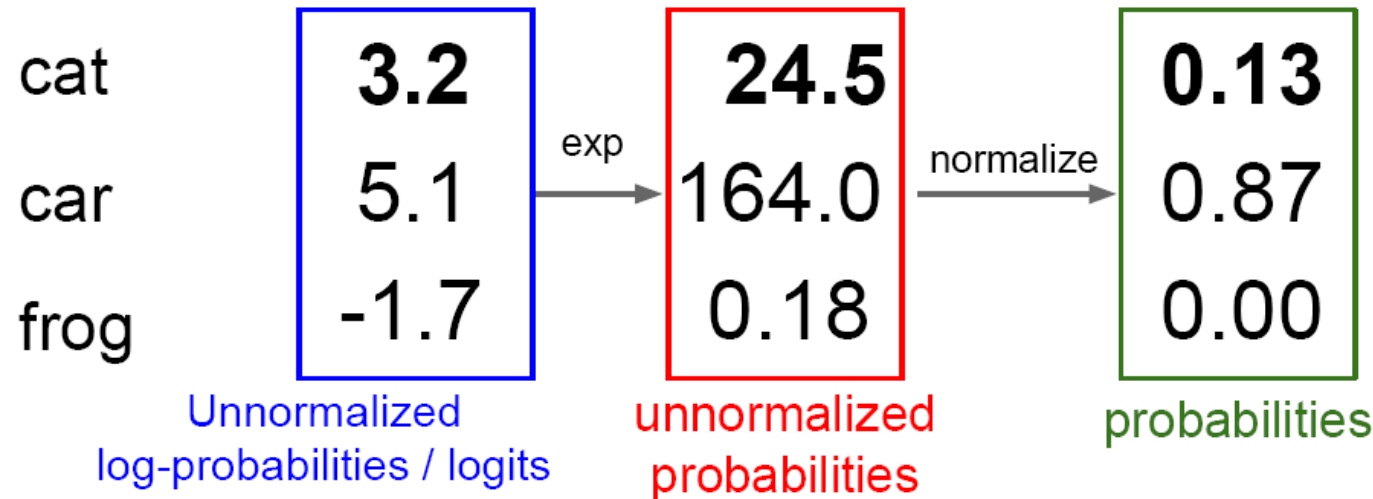


$$s = f(x_i; W)$$


$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$
 Softmax Function

Probabilities
must be ≥ 0

Probabilities
must sum to 1



Loss Function – Cross-entropy Loss





$$C = -\frac{1}{n} \sum_x [y \ln a + (1 - y) \ln(1 - a)]$$

Cross entropy is always larger than entropy; encoding symbols according to the wrong distribution will always make us use more bits.


Regularization - Dropout

In each forward pass, randomly set some neurons to zero
Probability of dropping is a hyperparameter; 0.5 is common




Regularization - Dropout

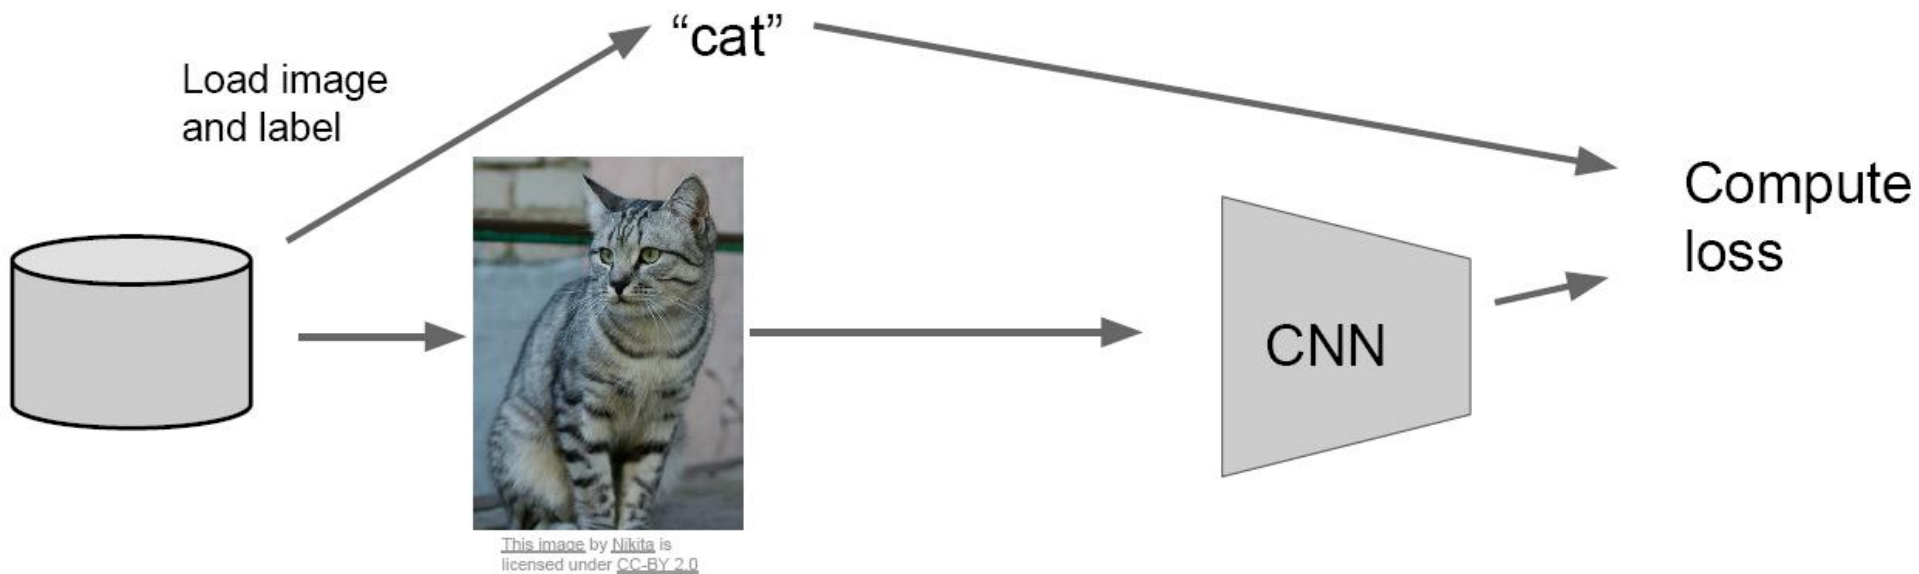
How can this possibly be a good idea?



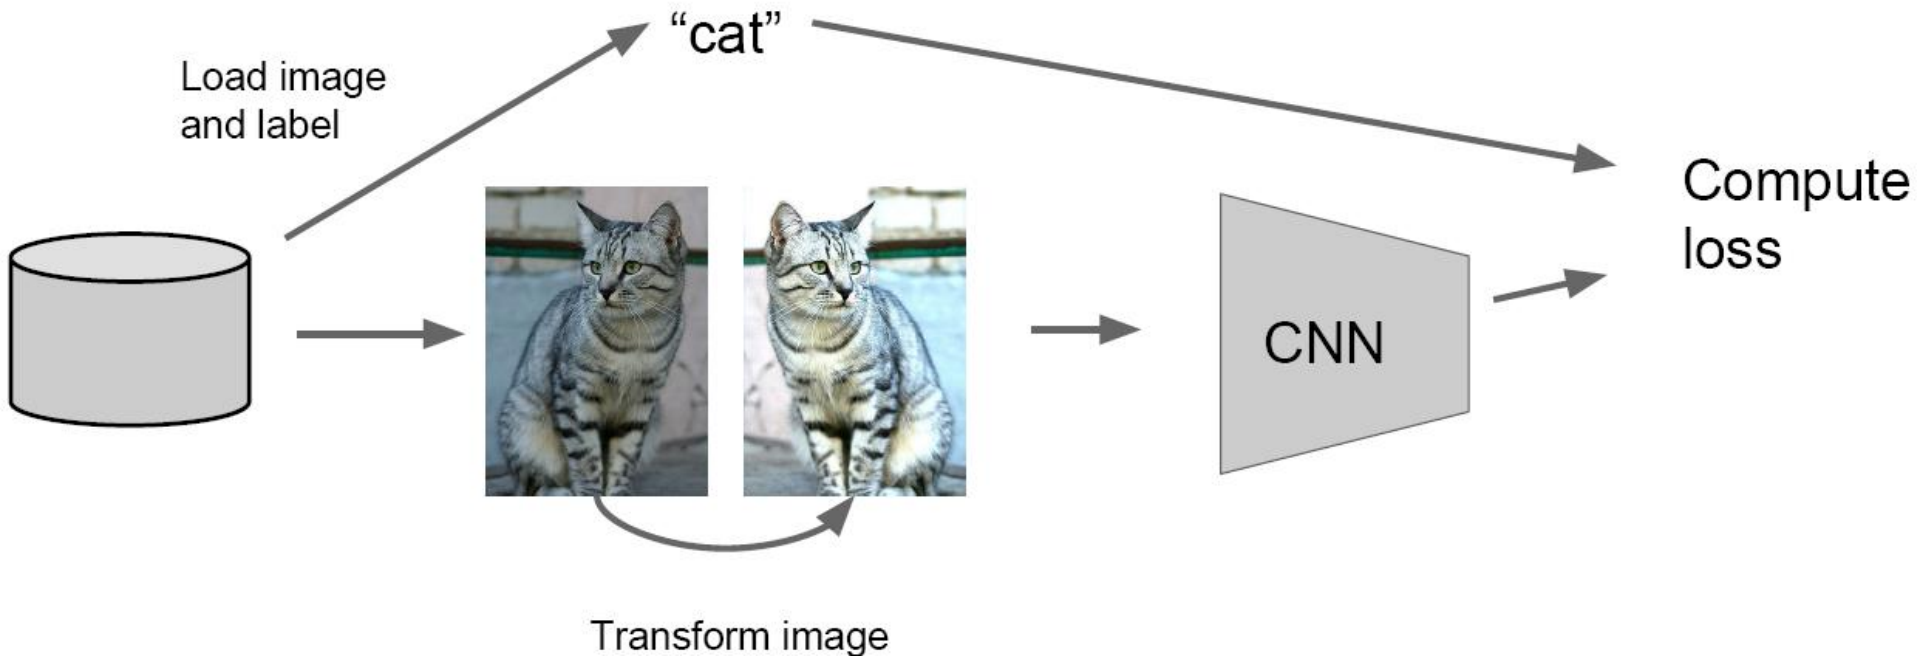
Forces the network to have a redundant representation;
Prevents co-adaptation of features




Regularization – Data Augmentation




Regularization – Data Augmentation



Regularization – Data Augmentation




Horizontal flips



Random crops and scale

Simple: Randomize
contrast and brightness



Color jitter

Random mix/combinations of :

- translation
- rotation
- stretching
- shearing,
- lens distortions, ... (go crazy)

Stochastic Gradient Descent

In stochastic (or "on-line") gradient descent, the true gradient is approximated by a gradient at a single example.

- Choose an initial vector of parameters w and learning rate η .
- Repeat until an approximate minimum is obtained:
 - Randomly shuffle examples in the training set.
 - For $i = 1, 2, \dots, n$, do:
 - $w := w - \eta \nabla Q_i(w)$.

Gradient descent: use all examples in each iteration


Stochastic gradient descent: use 1 example in each iteration

Mini-batch gradient descent: use b examples in each iteration

Representative CNN Networks

- LeNet-5
- AlexNet
- VGG
- Autoencoder
- ResNet
- GAN

LeNet-5



LeNet-5


1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

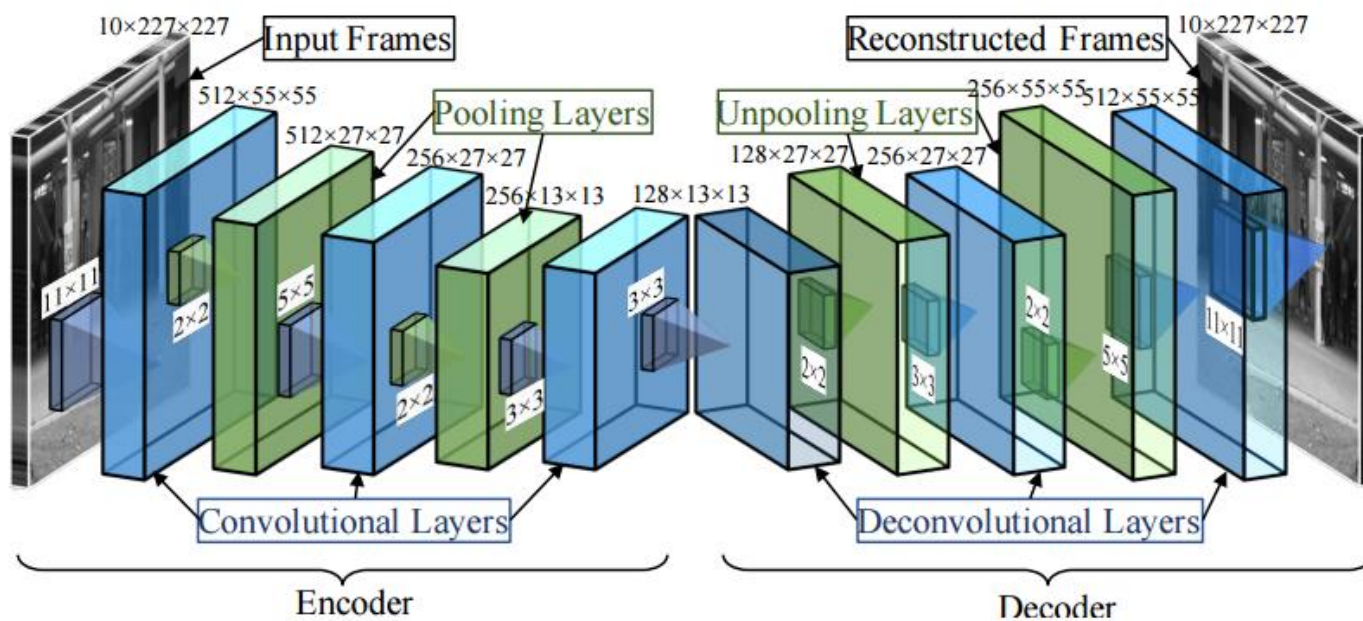

Convolution



Pooling


VGG19

Autoencoder



ResNet

Solution: Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



Use layers to
fit residual
 $F(x) = H(x) - x$
instead of
 $H(x)$ directly


Generative Adversarial Network

Real faces



Discriminator

Deep Convolutional Network (DCN)




Fake

Real


Generator

Deconvolutional Network (DN)


Random noise




Generated faces




Fluorescence Microscopy



U2OS



Actin



Endo

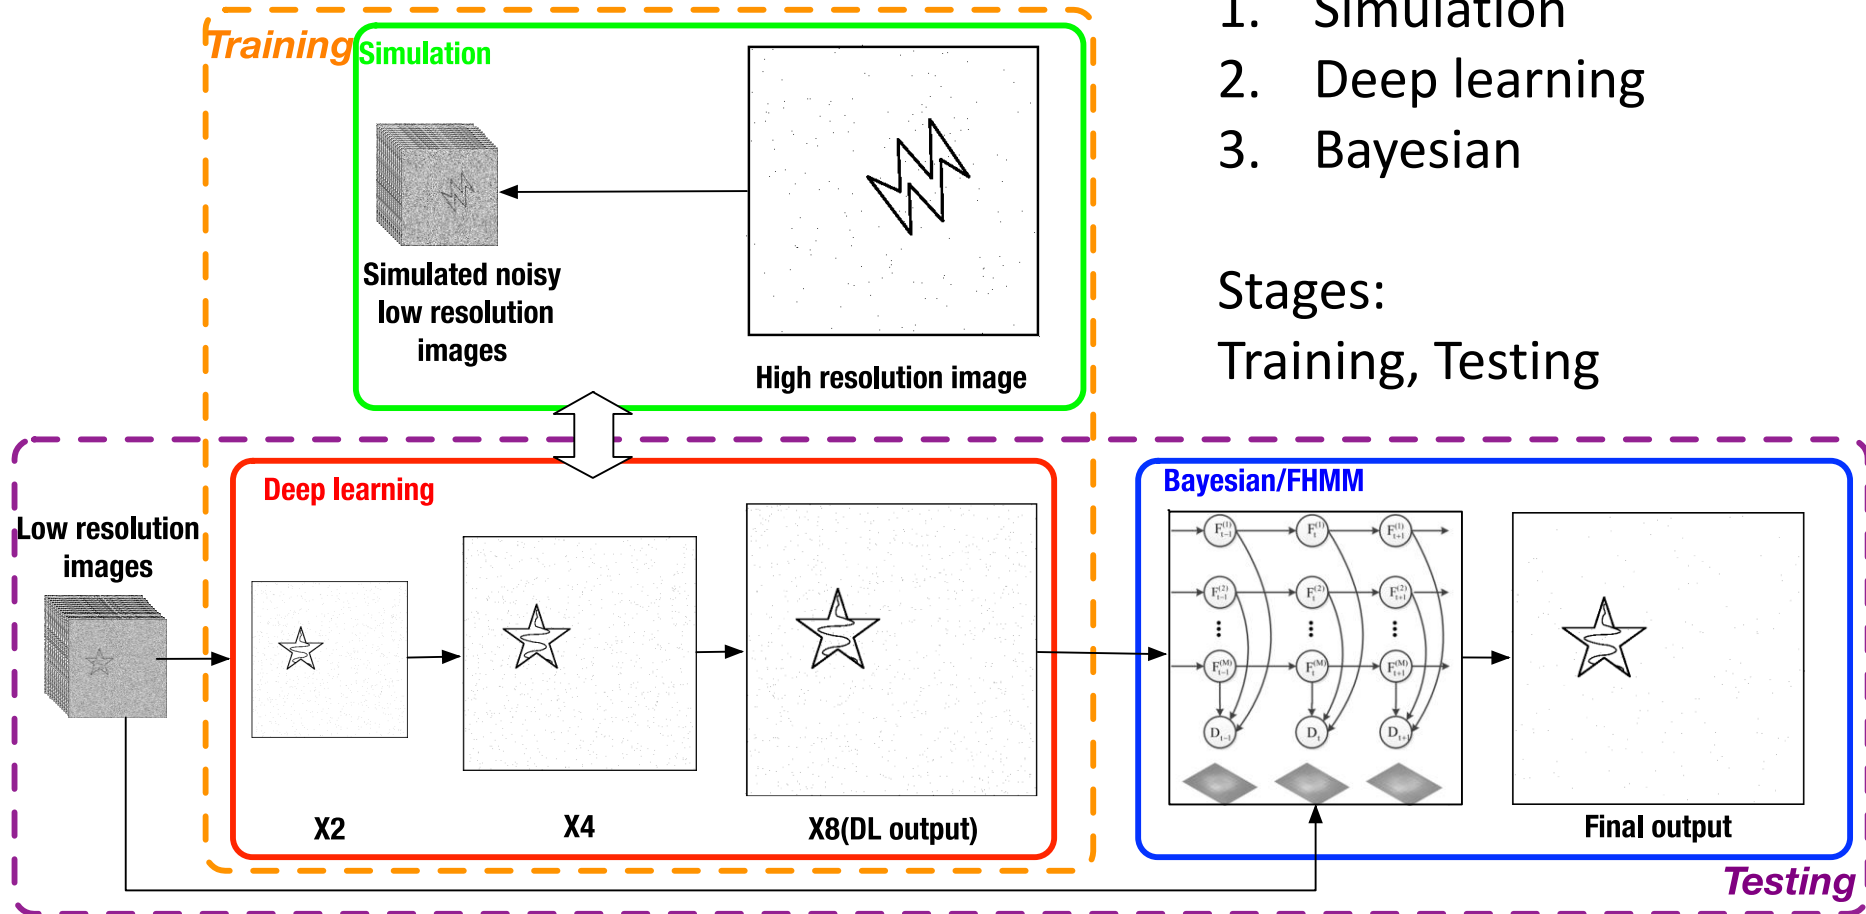
Overview of DLBI

Components:

1. Simulation
2. Deep learning
3. Bayesian

Stages:

Training, Testing






Performance on Real Data




U2OS

Actin




Endo



LR images




3B*



DLBI

Runtime




Compared to 3B:

DNN: **1500X** speed up


DLBI: **150X** speed up

- Large field reconstruction
- Real-time reconstruction


Large-field Reconstruction



Actin in U2OS (249*395)



DLBI: 200 frames (2K*3.2K)



PALM*: 20,000 frames

Backpropagation – Example

Backpropagation: a simple example


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial x}$$


Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Upstream
gradient

Local
gradient

Backpropagation – Example



Training ANN

- $\frac{\partial E}{\partial w_j^i} = \frac{\partial}{\partial w_j^i} \left[\frac{1}{2} \sum_{k=1}^p (y_k - c_k)^2 \right] = \frac{\partial}{\partial w_j^i} \left[\frac{1}{2} (y_j - c_j)^2 \right]$. That is, only the term where $k = j$ do we have any contribution made by w_j^i
- Recall $y_j = \sigma(s_j)$ and $s_j = \sum_{i=0}^n w_j^i x_i$, thus
$$\begin{aligned} \frac{\partial E}{\partial w_j^i} &= \frac{\partial}{\partial w_j^i} \left[\frac{1}{2} (y_j - c_j)^2 \right] = (y_j - c_j) \frac{\partial y_j}{\partial w_j^i} \\ &= (y_j - c_j) \frac{\partial y_j}{\partial s_j} \cdot \frac{\partial s_j}{\partial w_j^i} = (y_j - c_j) \frac{\partial y_j}{\partial s_j} \cdot \frac{\partial s_j}{\partial w_j^i} \\ &= (y_j - c_j) y_j (1 - y_j) x_i \end{aligned}$$

Training ANN


- $\frac{\partial E}{\partial w_j^i} = (y_j - c_j) y_j (1 - y_j) x_i$
- Define $\delta_j = (y_j - c_j) y_j (1 - y_j)$
- Thus $\frac{\partial E}{\partial w_j^i} = \delta_j x_i$
- More generally, $\delta_j = (y_j - c_j) \sigma'(s_j)$

Training ANN

- Now how do we use $\frac{\partial E}{\partial w_j^i}$?


– It is the gradient!

$$\frac{\partial E}{\partial w} = \left[\frac{\partial E}{\partial w^0}, \frac{\partial E}{\partial w_j^1}, \dots, \frac{\partial E}{\partial w_j^n} \right]$$



Training ANN

- Training rule: $\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$
- A small η means slow convergence, a big η means risks of jumping over global minimum
- Why “-”?
 - $\frac{\partial E}{\partial w_i}$ positive means Δw_i should be negative



Back-propagation Algorithm

- Initialize all weights to small random numbers.
- Until satisfied, Do
 - For each training example, Do
 - Input the training example to the network and compute the network outputs
 - For each output unit k : $\delta_k \leftarrow (z_k - C_k) z_k (1 - z_k)$
 - For each hidden unit h :
$$\delta_h \leftarrow z_h (1 - z_h) \sum_k \delta_k w_k^h$$
 - Update each network weight w_j^i : $w_j^i \leftarrow w_j^i + \Delta w_j^i$
Where $\Delta w_j^i = -\eta \delta_j x_i$